

Homework 1

Spring 2022

(Due: Friday, Jan 28, 2022, 4:59 pm Eastern Time)

Please submit your homework through **gradescope**. You can write, scan, type, etc. But for the convenience of grading, please merge everything into a **single PDF**.

Objective

To complete this homework assignment, you need to read the Background chapter of the lecture note, available on the course website under the page **note**. After this homework, you will be familiar with:

- (a) Basic plotting tools in Python, for both 1D and 2D plots. Drawing random samples in Python.
- (b) Cross-validation. See my book *Introduction to Probability for Data Science*, Chapter 3.2.5.
- (c) Gaussian whitening. See my book Chapter 5.7.4.
- (d) Basic ideas of regression. See my book Chapter 7.1.4.

You will be asked some of these questions in Quiz 1.

Exercise 1: Histogram and Cross-Validation

Let X be a random variable with $X \sim \mathcal{N}(\mu, \sigma^2)$. The PDF of X is written explicitly as

$$f_X(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}. \quad (1)$$

- (a) Let $\mu = 0$ and $\sigma = 1$ so that $X \sim \mathcal{N}(0, 1)$. Plot $f_X(x)$ using `matplotlib.pyplot.plot` for the range $x \in [-3, 3]$. Use `matplotlib.pyplot.savefig` to save your figure.
- (b) Let us investigate the use of histograms in data visualization.
 - (i) Use `numpy.random.normal` to draw 1000 random samples from $\mathcal{N}(0, 1)$.
 - (ii) Make two histogram plots using `matplotlib.pyplot.hist`, with the number of bins m set to 4 and 1000.
 - (iii) Use `scipy.stats.norm.fit` to estimate the mean and standard deviation of your data. Report the estimated values.
 - (iv) Plot the fitted gaussian curve on top of the two histogram plots using `scipy.stats.norm.pdf`.
 - (v) (Optional) Ask yourself the following questions: Are the two histograms representative of your data's distribution? How are they different in terms of data representation?
- (c) A practical way to estimate the optimal bin width is to make use of what is called the **cross validation estimator of risk** of the dataset. Denoting $h = (\max \text{ data value} - \min \text{ data value})/m$ as the bin width, with $m =$ the number of bins (assuming you applied no rescaling to your raw data), we seek h^* that minimizes the risk $\hat{J}(h)$, expressed as follows:

$$\hat{J}(h) = \frac{2}{h(n-1)} - \frac{n+1}{h(n-1)} \sum_{j=1}^m \hat{p}_j^2, \quad (2)$$

where $\{\hat{p}_j\}_{j=1}^m$ is the empirical probability of a sample falling into each bin, and n is the total number of samples.

- (i) Plot $\hat{J}(h)$ with respect to m the number of bins, for $m = 1, 2, \dots, 200$.
- (ii) Find the m^* that minimizes $\hat{J}(h)$, plot the histogram of your data with that m^* .
- (iii) Plot the Gaussian curve fitted to your data on top of your histogram.

Note: For additional discussions about this cross-validation technique, visit my book *Introduction to Probability for Data Science*, Chapter 3.2.5. More advanced materials can be found in the supplementary note of this homework.

Exercise 2: Gaussian Whitening

In this exercise, we consider the following question: suppose that we are given a random number generator that can only generate zero-mean unit variance Gaussians, i.e., $\mathbf{X} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, how do we transform the distribution of \mathbf{X} to an arbitrary Gaussian distribution? We will first derive a few equations, and then verify them with an empirical example, by drawing samples from the 2D Gaussian, applying the transform to the dataset, and checking if the transformed dataset really takes the form of the desired Gaussian.

- (a) Let $\mathbf{X} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ be a 2D Gaussian. The PDF of \mathbf{X} is given by

$$f_{\mathbf{X}}(\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^2 |\boldsymbol{\Sigma}|}} \exp \left\{ -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right\}, \quad (3)$$

where in this exercise we assume

$$\mathbf{X} = \begin{bmatrix} X_1 \\ X_2 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}, \quad \boldsymbol{\mu} = \begin{bmatrix} 2 \\ 6 \end{bmatrix}, \quad \text{and} \quad \boldsymbol{\Sigma} = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix} \quad (4)$$

- (i) Simplify the expression $f_{\mathbf{X}}(\mathbf{x})$ for the particular choices of $\boldsymbol{\mu}$ and $\boldsymbol{\Sigma}$ here. Show your derivation.
 - (ii) Using `matplotlib.pyplot.contour`, plot the contour of $f_{\mathbf{X}}(\mathbf{x})$ for the range $\mathbf{x} \in [-1, 5] \times [0, 10]$.
- (b) Suppose $\mathbf{X} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. We would like to derive a transformation that can map \mathbf{X} to an arbitrary Gaussian.

- (i) Let $\mathbf{X} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ be a d -dimensional random vector. Let $\mathbf{A} \in \mathbb{R}^{d \times d}$ and $\mathbf{b} \in \mathbb{R}^d$. Let $\mathbf{Y} = \mathbf{A}\mathbf{X} + \mathbf{b}$ be an affine transformation of \mathbf{X} . Let $\boldsymbol{\mu}_{\mathbf{Y}} \stackrel{\text{def}}{=} \mathbb{E}[\mathbf{Y}]$ be the mean vector and $\boldsymbol{\Sigma}_{\mathbf{Y}} \stackrel{\text{def}}{=} \mathbb{E}[(\mathbf{Y} - \boldsymbol{\mu}_{\mathbf{Y}})(\mathbf{Y} - \boldsymbol{\mu}_{\mathbf{Y}})^T]$ be the covariance matrix. Show that

$$\boldsymbol{\mu}_{\mathbf{Y}} = \mathbf{b}, \quad \text{and} \quad \boldsymbol{\Sigma}_{\mathbf{Y}} = \mathbf{A}\mathbf{A}^T. \quad (5)$$

- (ii) Show that $\boldsymbol{\Sigma}_{\mathbf{Y}}$ is symmetric positive semi-definite.
- (iii)** Under what condition on \mathbf{A} would $\boldsymbol{\Sigma}_{\mathbf{Y}}$ become a symmetric positive definite matrix?
- (iv) Consider a random variable $\mathbf{Y} \sim \mathcal{N}(\boldsymbol{\mu}_{\mathbf{Y}}, \boldsymbol{\Sigma}_{\mathbf{Y}})$ such that

$$\boldsymbol{\mu}_{\mathbf{Y}} = \begin{bmatrix} 2 \\ 6 \end{bmatrix}, \quad \text{and} \quad \boldsymbol{\Sigma}_{\mathbf{Y}} = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}.$$

Determine \mathbf{A} and \mathbf{b} which could satisfy Equation (5).

Hint: Consider eigen-decomposition of $\boldsymbol{\Sigma}_{\mathbf{Y}}$. You may compute the eigen-decomposition numerically.

- (c) Now let us verify our results from part (b) with an empirical example.

- (i) Use `numpy.random.multivariate_normal` to draw 5000 random samples from the 2D standard normal distribution, and make a scatter plot of the data point using `matplotlib.pyplot.scatter`.
- (ii) Apply the affine transformation you derived in part (b)(iv) to the data points, and make a scatter plot of the transformed data points. **Now check your answer by using the Python function `numpy.linalg.eig` to obtain the trasformation and making a new scatter plot of the transformed data points.**
- (iii) (Optional) Do your results from parts (c)(i) and (ii) support your theoretical findings from part (b)?

You can find more information about Gaussian whitening in my book *Introduction to Probability for Data Science*, Chapter 5.7.4.

Exercise 3: Linear Regression

Let us consider a polynomial fitting problem. We assume the following model:

$$y = \beta_0 + \beta_1 L_1(x) + \beta_2 L_2(x) + \dots + \beta_p L_p(x) + \epsilon, \quad (6)$$

where $L_p(x)$ is the p -th Legendre polynomial, β_j are the coefficients, and ϵ is the error term. In Python, if you have specified a list of values of x , evaluating the Legendre polynomial is quite straight forward:

```
import numpy as np
from scipy.special import eval_legendre
x = np.linspace(-1,1,50) # 50 points in the interval [-1,1]
L4 = eval_legendre(4,x)   # evaluate the 4th order Legendre polynomial for x
```

- (a) Let $\beta_0 = -0.001$, $\beta_1 = 0.01$, $\beta_2 = 0.55$, $\beta_3 = 1.5$, $\beta_4 = 1.2$, and let $\epsilon \sim \text{Gaussian}(0, 0.2^2)$. Generate 50 points of y over the interval $x = \text{np.linspace}(-1,1,50)$. That is,

```
x = np.linspace(-1,1,50) # 50 points in the interval [-1,1]
y = ... # fill this line
```

Scatter plot the data.

- (b) Given the $N = 50$ data points, formulate the linear regression problem. Specifically, write down the expression

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\beta\|^2. \quad (7)$$

What are \mathbf{y} , \mathbf{X} , and β ? Derive the optimal solution for this simple regression problem. Express your answer in terms of \mathbf{X} and \mathbf{y} .

- (c) Write a Python code to compute the solution. Overlay your predicted curve with the scattered plot. For solving the regression problem, you can call `numpy.linalg.lstsq`.
- (d) For the \mathbf{y} you have generated, make 5 outlier points using the code below:

```
# ...
idx = [10,16,23,37,45]   # these are the locations of the outliers
y[idx] = 5               # set the outliers to have a value 5
# ...
```

Run your code in (c) again. Comment on the difference.

(e) Consider the optimization

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\beta\|_1. \quad (8)$$

Convert the problem into a linear programming problem. Express your solution in the linear programming form:

$$\begin{aligned} & \underset{\mathbf{x}}{\operatorname{minimize}} && \mathbf{c}^T \mathbf{x} \\ & \text{subject to} && \mathbf{A}\mathbf{x} \leq \mathbf{b}. \end{aligned} \quad (9)$$

What are \mathbf{c} , \mathbf{x} , \mathbf{A} , and \mathbf{b} ?

(f) Solve the linear programming problem in Python using `scipy.optimize.linprog`, for the corrupted data in (d). Scatter plot the data, and overlay with your predicted curve. Hint: Remember to set `bounds=(None, None)` when you call `scipy.optimize.linprog`, because the variables in `linprog` are non-negative by default.

For this problem, you may want to check my book *Introduction to Probability for Data Science*, Chapter 7.1.4.

Exercise 4: Project: Check Point 1

By now I believe you should have read the course project instructions. If not, please read them now. The objective of this series of “check points” is to keep track of your progress, so that you will not wait until the last minute and then become panic. For check point 1, I want you to complete the following tasks:

1. **Latex.** I only accept final reports typed in Latex, using the ICML 2021 template. You can use overleaf (a free online platform) to type your report. Download the ICML template, put it in overleaf. Change the title to your project title, and change the author name to your name. (By default the ICML template uses the review mode. Please switch it to the camera ready mode.)

When typing your name, please tell us: Your name, your major (e.g., ECE, AAE, etc), and your level (e.g., Online MS, PhD, Undergrad, etc)

2. **Topic.** There are four topics listed in the course website. Read them carefully. Think about which one you want to work on. Put a tentative title to your latex file. Your title has to be informative about what *you* want to do. E.g., “Comparing Baseline Methods for Noise2Noise”, “Re-Implementation of One-Size-Fits-All”, etc.
3. **Datasets.** Every project will use some kind of datasets. Go to Google, search around, and find the most suitable datasets. List at least two possible datasets.
4. **Codes.** Some projects have existing codes on the internet. Find them out, and list them.
5. **References.** List at least one main references for your project, and at least 5 additional references. Follow the ICML instructions about the format of the .bib file.

Therefore, in this check point, I am expecting you to attach a one page PDF compiled from LaTeX (in ICML 2021 template). The one page should contain a title, your name, a one sentence summary of the project that you want to work on, a list of datasets you have found on the internet, the available codes you have found, and a list of references.

I understand that many of you have not yet decided which project to work on. You are free to change the project topic any time during the semester. You are also free to change the dataset, the code, the references, etc. As I said, the purpose here is to help you get started. Even if you have not yet decided, you should still do this exercise, because it will help you get familiar with the steps.